

Advertising Effects

Ben Heim, Marcell Milo-Sidlo, Pierre Chan, Julia Luo, Alicia Soosai

Contents

1	Overview	2
2	Data	2
2.1	Brands and product modules	2
3	3 Data preparation	3
3.1	3.1 RMS scanner data (<code>move</code>)	3
3.2	3.2 Ad Intel advertising data (<code>adv_DT</code>)	5
3.3	3.3 Calculate adstock/goodwill	8
3.4	Merge scanner and advertising data	10
3.5	Reshape the data	10
3.6	Time fixed effects	11
4	Data inspection	12
4.1	Time-series of advertising levels	12
4.2	Overall advertising variation	13
5	Advertising effect estimation	15

```
library(bit64)
library(data.table)
library(RcppRoll)
library(ggplot2)
library(fixest)
library(knitr)
```

1 Overview

In this assignment we estimate the causal short and long-run effects of advertising on demand. The assignment is closely related to the paper “TV Advertising Effectiveness and Profitability: Generalizable Results from 288 Brands” (2001, *Econometrica*) by Shapiro, Hitsch, and Tuchman.

We first combine store-level sales data from the Nielsen RMS scanner data set with DMA-level advertising exposure data from the Nielsen Ad Intel advertising data set. We then estimate ad effects based on a within-market strategy controlling for cross-sectional heterogeneity across markets.

2 Data

2.1 Brands and product modules

Data location:

```
data_folder = "data"
```

The table `brands_DT` in the file `Brands_a3.RData` provides information on the available product categories (product modules) and brands, including the “focal” brands for which we may estimate advertising effects.

```
brands = load(paste0(data_folder, "/Brands_a3.RData"))
```

product_module_code	product_module_desc	brand_code_uc	brand_descr	focal_brand
1484	SOFT DRINKS - CARBONATED	531429	COCA-COLA R	TRUE
8412	ANTACIDS	621727	PRILOSEC	TRUE
1553	SOFT DRINKS - LOW CALORIE	531433	COCA-COLA ZERO DT	TRUE

Choose Prilosec in the Antacids category for your analysis. Later, you can **optionally** repeat your analysis for the other brands.

```
selected_module = 8412
selected_brand = 621727
```

3 3 Data preparation

To prepare and build the data for the main analysis, load the brand and store meta-data in `Brands_a3.RData` and `stores_dma.RData`, the RMS store-level scanner (movement) data, and the Nielsen Ad Intel DMA-level TV advertising data. The scanner data and advertising data are named according to the product module, such as `move_8412.RData` and `adv_8412.RData`.

Both the RMS scanner data and the Ad Intel advertising data include information for the top four brands in the category (product module). To make our analysis computationally more manageable we will not distinguish among all individual competing brands, but instead we will aggregate all competitors into one single brand.

```
brands = load(paste0(data_folder, "/Brands_a3.RData"))
stores = load(paste0(data_folder, "/stores_dma.RData"))
move_data = load(paste0(data_folder, "/move_8412.RData"))
advertising_data = load(paste0(data_folder, "/adv_8412.RData"))
```

3.1 3.1 RMS scanner data (move)

Let us start manipulating the ‘move’ dataset.

For consistency, rename the `units` to `quantity` and `promo_percentage` to `promotion` (use the `setnames` command). The promotion variable captures promotional activity as a continuous variable with values between 0 and 1.

Create the variable `brand_name` that we will use to distinguish between the own and aggregate competitor variables. The brand name `own` corresponds to the focal brand (Prilosec in our case), and `comp` (or any other name that you prefer) corresponds to the aggregate competitor brand.

We need to aggregate the data for each store/week observation, separately for the `own` and `comp` data. To aggregate prices and promotions we can take the simple arithmetic `mean` over all competitor brands (a weighted mean may be preferable but is not necessary in this analysis where prices and promotions largely serve as controls, not as the main marketing mix variables of interest). Aggregate quantities can be obtained as the `sum` over brand-level quantities.

```
#Step 1: Renaming (use set names)
#Step 1a: Units -> Quantity
#Step 1b: promo_percentage -> promotion
library(data.table)
setnames(move, old = "units", new = "quantity")
setnames(move, old = "promo_percentage", new = "promotion")

head(move)
```

Key: <brand_code_uc, store_code_uc, week_end>

	brand_code_uc	store_code_uc	week_end	quantity	price	promotion
	<int>	<int>	<Date>	<num>	<num>	<num>
1:	536746	2324	2010-01-02	4520	0.2842310	0
2:	536746	2324	2010-01-09	5456	0.2864958	0
3:	536746	2324	2010-01-16	5102	0.2864958	0
4:	536746	2324	2010-01-23	4950	0.2864041	0
5:	536746	2324	2010-01-30	4697	0.2861501	0
6:	536746	2324	2010-02-06	6090	0.2863082	0

```
#Step 2: Identifying 'own' (Prilosec) vs. 'comp' (non-Prilosec)
move$brand_name = ifelse(move$brand_code_uc == selected_brand, "own", "comp")
colnames(move)
```

```
[1] "brand_code_uc" "store_code_uc" "week_end"      "quantity"
[5] "price"         "promotion"     "brand_name"
```

```
head(move)
```

```
Key: <brand_code_uc, store_code_uc, week_end>
```

	brand_code_uc	store_code_uc	week_end	quantity	price	promotion
	<int>	<int>	<Date>	<num>	<num>	<num>
1:	536746	2324	2010-01-02	4520	0.2842310	0
2:	536746	2324	2010-01-09	5456	0.2864958	0
3:	536746	2324	2010-01-16	5102	0.2864958	0
4:	536746	2324	2010-01-23	4950	0.2864041	0
5:	536746	2324	2010-01-30	4697	0.2861501	0
6:	536746	2324	2010-02-06	6090	0.2863082	0

	brand_name
	<char>
1:	comp
2:	comp
3:	comp
4:	comp
5:	comp
6:	comp

```
#Step 3: Aggregating the data for each store/observation
```

```
#For 'comp'
```

```
move <- move[, .(
  price = mean(price, na.rm = TRUE),      # Mean price
  promotion = mean(promotion, na.rm = TRUE), # Mean promotion
  quantity = sum(quantity, na.rm = TRUE)   # Sum quantities
), by = .(store_code_uc, week_end, brand_name)]
```

```
head(move)
```

	store_code_uc	week_end	brand_name	price	promotion	quantity
	<int>	<Date>	<char>	<num>	<num>	<num>
1:	2324	2010-01-02	comp	0.4396163	0.167566	5048
2:	2324	2010-01-09	comp	0.4424408	0.000000	6008
3:	2324	2010-01-16	comp	0.4399273	0.000000	5898
4:	2324	2010-01-23	comp	0.4367092	0.167566	5508
5:	2324	2010-01-30	comp	0.4327154	0.167566	5375
6:	2324	2010-02-06	comp	0.4339731	0.167566	6310

Later, when we merge the RMS scanner data with the Ad Intel advertising data, we need a common key between the two data sets. This key will be provided by the DMA code and the date. Hence, we need to merge the `dma_code` found in the `stores` table with the RMS movement data.

Now merge the `dma_code` with the movement data.

```
#Merge stores_dma w/move
```

```
move = merge(move, stores_dma, by = "store_code_uc")
```

```
head(move)
```

```
Key: <store_code_uc>
```

	store_code_uc	week_end	brand_name	price	promotion	quantity	dma_code
	<int>	<Date>	<char>	<num>	<num>	<num>	<int>
1:	2324	2010-01-02	comp	0.4396163	0.167566	5048	602
2:	2324	2010-01-09	comp	0.4424408	0.000000	6008	602

3:	2324	2010-01-16	comp	0.4399273	0.000000	5898	602
4:	2324	2010-01-23	comp	0.4367092	0.167566	5508	602
5:	2324	2010-01-30	comp	0.4327154	0.167566	5375	602
6:	2324	2010-02-06	comp	0.4339731	0.167566	6310	602

```

dma_descr
<char>
1: CHICAGO IL
2: CHICAGO IL
3: CHICAGO IL
4: CHICAGO IL
5: CHICAGO IL
6: CHICAGO IL

```

3.2 Ad Intel advertising data (adv_DT)

The table `adv_DT` contains information on brand-level GRPs (gross rating points) for each DMA/week combination. The original data are more disaggregated, and include individual occurrences on a specific date and at a specific time and the corresponding number of impressions. `adv_DT` is based on the original data, aggregated at the DMA/week level.

Weeks are indicated by `week_end`, where the corresponding date is always a Saturday. We use Saturdays so that the `week_end` variable in the advertising data corresponds to the date convention in the RMS scanner data, where `week_end` also corresponds to a Saturday.

The data contain two variables to measure brand-level GRPs, `grp_direct` and `grp_indirect`. `grp_direct` records GRPs for which we can create a direct, unambiguous match between the brand name in the scanner data and the name of the advertised brand. Sometimes, however, it is not entirely clear if we should associate an ad in the Ad Intel data with the brand in the RMS data. For example, should we count ads for BUD LIGHT BEER LIME when measuring the GRPs that might affect sales of BUD LIGHT BEER? As such matches are somewhat debatable, we record the corresponding GRPs in the variable `grp_indirect`.

The data do not contain observations for all DMA/week combinations during the observation period. In particular, no DMA/week record is included if there was no corresponding advertising activity. For our purposes, however, it is important to capture that the number of GRPs was 0 for such observations. Hence, we need to “fill the gaps” in the data set.

`data.table` makes it easy to achieve this goal. Let’s illustrate using a simple example:

```

set.seed(444)
DT = data.table(dma = rep(LETTERS[1:2], each = 5),
               week = 1:5,
               x = round(runif(10, min = 0, max = 20)))
DT = DT[-c(2, 5, 9)]
DT

```

	dma	week	x
	<char>	<int>	<num>
1:	A	1	3
2:	A	3	8
3:	A	4	7
4:	B	1	12
5:	B	2	11
6:	B	3	1
7:	B	5	6

In `DT`, the observations for weeks 2 and 5 in market A and week 4 in market B are missing.

To fill the holes, we need to key the `data.table` to specify the dimensions—here the `dma` and `week`. Then we perform a *cross join* using `CJ` (see `?CJ`). In particular, for each of the variables along which `DT` is keyed we specify the full set of values that the final `data.table` should contain. In this example, we want to include the markets A and B and all weeks, 1-5.

```
setkey(DT, dma, week)
DT = DT[CJ(c("A", "B"), 1:5)]
DT
```

```
Key: <dma, week>
   dma week    x
   <char> <int> <num>
1:     A     1     3
2:     A     2    NA
3:     A     3     8
4:     A     4     7
5:     A     5    NA
6:     B     1    12
7:     B     2    11
8:     B     3     1
9:     B     4    NA
10:    B     5     6
```

We can replace all missing values (NA) with another value, say -111, like this:

```
DT[is.na(DT)] = -111
DT
```

```
Key: <dma, week>
   dma week    x
   <char> <int> <num>
1:     A     1     3
2:     A     2  -111
3:     A     3     8
4:     A     4     7
5:     A     5  -111
6:     B     1    12
7:     B     2    11
8:     B     3     1
9:     B     4  -111
10:    B     5     6
```

Use this technique to expand the advertising data in `adv_DT`, using a cross join along all `brands`, `dma_codes`, and `weeks`:

```
brands      = unique(adv_DT$brand_code_uc)
dma_codes   = unique(adv_DT$dma_code)
weeks       = seq(from = min(adv_DT$week_end), to = max(adv_DT$week_end),
                  by = "week")
```

Now perform the cross join and set missing values to 0.

```
#Cross Join
library(data.table)
setDT(adv_DT)
complete_set = CJ(brand_code_uc = brands, dma_code = dma_codes,
                  week_end = weeks)
```

```
adv_DT = complete_set[adv_DT, on = .(brand_code_uc, dma_code, week_end)]
adv_DT[is.na(adv_DT)] = 0
```

```
adv_DT
```

```
Key: <brand_code_uc, dma_code, week_end>
```

	brand_code_uc	dma_code	week_end	grp_direct	grp_indirect
	<int>	<num>	<Date>	<num>	<num>
1:	616173	500	2010-01-02	2.130524	0
2:	616173	500	2010-01-09	71.706024	0
3:	616173	500	2010-01-16	78.969427	0
4:	616173	500	2010-01-23	98.092552	0
5:	616173	500	2010-01-30	69.769588	0

64940:	621727	881	2014-12-06	121.749822	0
64941:	621727	881	2014-12-13	154.261782	0
64942:	621727	881	2014-12-20	155.543606	0
64943:	621727	881	2014-12-27	246.334964	0
64944:	621727	881	2015-01-03	136.245386	0

Create own and competitor names, and then aggregate the data at the DMA/week level, similar to what we did with the RMS scanner data. In particular, aggregate based on the sum of GRPs (separately for `grp_direct` and `grp_indirect`).

```
#Locating own and comps within advertising datatable
```

```
selected_brand = 621727
```

```
adv_DT[, brand_name := ifelse(adv_DT$brand_code_uc == selected_brand,
                              "own", "comp")]
```

```
adv_DT <- adv_DT[, .(
  grp_direct = sum(grp_direct, na.rm = TRUE),      # Sum of direct GRPs
  grp_indirect = sum(grp_indirect, na.rm = TRUE)   # Sum of indirect GRPs
), by = .(dma_code, week_end, brand_name)]
```

At this stage we need to decide if we want to measure GRPs using only `grp_direct` or also including `grp_indirect`. I propose to take the broader measure, and sum the GRPs from the two variables to create a combined `grp` measure. You can later check if your results are robust if you use `grp_direct` only (this robustness analysis is optional).

```
#Direct & Indirect GRPs
```

```
adv_DT$grp = adv_DT$grp_direct + adv_DT$grp_indirect
```

Note: In the Antacids category, `grp_indirect` only contains the value 0 and is therefore not relevant. However, if you work with the data in the other categories, `grp_indirect` contains non-zero values.

3.3 Calculate adstock/goodwill

Advertising is likely to have long-run effects on demand. Hence, we will calculate adstock or goodwill variables for own and competitor advertising. We will use the following, widely-used adstock specification (a_t is advertising in period t):

$$g_t = \sum_{l=0}^L \delta^l \log(1 + a_{t-l}) = \log(1 + a_t) + \delta \log(1 + a_{t-1}) + \cdots + \delta^L \log(1 + a_{t-L})$$

We add 1 to the advertising levels (GRPs) before taking the log to deal with the large number of zeros in the GRP data.

Here is a particularly easy and fast approach to calculate adstocks. First, define the adstock parameters—the number of lags and the carry-over factor δ .

```
N_lags = 52
delta = 0.7
```

Then calculate the geometric weights based on the carry-over factor.

```
geom_weights = cumprod(c(1.0, rep(delta, times = N_lags)))
geom_weights = sort(geom_weights)
tail(geom_weights)
```

```
[1] 0.16807 0.24010 0.34300 0.49000 0.70000 1.00000
```

Now we can calculate the adstock variable using the `roll_sum` function in the `RcppRoll` package.

```
adv_DT[, grp := grp_direct + grp_indirect]

setkey(adv_DT, brand_name, dma_code, week_end)

adv_DT[, adstock := roll_sum(log(1 + grp),
                             n = N_lags + 1,
                             weights = geom_weights,
                             normalize = FALSE,
                             align = "right", fill = NA),
        by = .(brand_name, dma_code)]
```

Explanations:

1. Key the table along the cross-sectional units (brand name and DMA), then along the time variable. This step is *crucial*! If the table is not correctly sorted, the time-series order of the advertising data will be incorrect.
2. Use the `roll_sum` function based on `log(1+grp)`. `n` indicates the total number of elements in the rolling sum, and `weights` indicates the weights for each element in the sum. `normalize = FALSE` tells the function to leave the `weights` untouched, `align = "right"` indicates to use all data above the current row in the data table to calculate the sum, and `fill = NA` indicates to fill in missing values for the first rows for which there are not enough elements to take the sum.

Alternatively, you could code your own weighted sum function:

```
weightedSum <- function(x, w) {
  T = length(x)
  L = length(w) - 1
  y = rep_len(NA, T)
  for (i in (L+1):T) y[i] = sum(x[(i-L):i]*w)
```



```

    return(y)
}

```

Let's compare the execution speed:

```

time_a = system.time(adv_DT[, stock_a := weightedSum(log(1+grp), geom_weights),
                                                         by = .(brand_name, dma_code)])
time_a

```

```

      user  system elapsed
0.351    0.018    0.094

```

```

time_b = system.time(adv_DT[, stock_b := roll_sum(log(1+grp), n = N_lags+1,
                                                         weights = geom_weights,
                                                         normalize = FALSE,
                                                         align = "right", fill = NA),
                                                         by = .(brand_name, dma_code)])
time_b

```

```

      user  system elapsed
0.005    0.000    0.005

```

Even though the `weightedSum` function is fast, the speed difference with respect to the optimized code in `RcppRoll` is large.

```

(time_a/time_b)[3]

```

```

elapsed
18.8

```

Lesson: Instead of reinventing the wheel, spend a few minutes searching the Internet to see if someone has already written a package that solves your coding problems.

3.4 Merge scanner and advertising data

Merge (join) the advertising data with the scanner data based on brand name, DMA code, and week.

```
head(move)
```

```
Key: <store_code_uc>
```

	store_code_uc	week_end	brand_name	price	promotion	quantity	dma_code
	<int>	<Date>	<char>	<num>	<num>	<num>	<int>
1:	2324	2010-01-02	comp	0.4396163	0.167566	5048	602
2:	2324	2010-01-09	comp	0.4424408	0.000000	6008	602
3:	2324	2010-01-16	comp	0.4399273	0.000000	5898	602
4:	2324	2010-01-23	comp	0.4367092	0.167566	5508	602
5:	2324	2010-01-30	comp	0.4327154	0.167566	5375	602
6:	2324	2010-02-06	comp	0.4339731	0.167566	6310	602

	dma_descr
	<char>
1:	CHICAGO IL
2:	CHICAGO IL
3:	CHICAGO IL
4:	CHICAGO IL
5:	CHICAGO IL
6:	CHICAGO IL

```
adv_DT[, stock_a := NULL]
adv_DT[, stock_b := NULL]
move <- merge(move, adv_DT, by = c("brand_name", "dma_code", "week_end"))
head(move)
```

```
Key: <brand_name, dma_code, week_end>
```

	brand_name	dma_code	week_end	store_code_uc	price	promotion	quantity
	<char>	<int>	<Date>	<int>	<num>	<num>	<num>
1:	comp	500	2010-01-02	88153	0.5957322	0.00000000	1432
2:	comp	500	2010-01-02	95752	0.5644275	0.08261606	1946
3:	comp	500	2010-01-02	123214	0.6173410	0.00000000	1592
4:	comp	500	2010-01-02	129685	0.5777817	0.20306278	1146
5:	comp	500	2010-01-02	189538	0.5948342	0.00000000	1698
6:	comp	500	2010-01-02	366762	0.4063640	0.23904801	4388

	dma_descr	grp_direct	grp_indirect	grp	adstock
	<char>	<num>	<num>	<num>	<num>
1:	PORTLAND-AUBURN ME	211.2821	0	211.2821	NA
2:	PORTLAND-AUBURN ME	211.2821	0	211.2821	NA
3:	PORTLAND-AUBURN ME	211.2821	0	211.2821	NA
4:	PORTLAND-AUBURN ME	211.2821	0	211.2821	NA
5:	PORTLAND-AUBURN ME	211.2821	0	211.2821	NA
6:	PORTLAND-AUBURN ME	211.2821	0	211.2821	NA

3.5 Reshape the data

Use `dcast` to reshape the data from long to wide format. The store code and week variable are the main row identifiers. Quantity, price, promotion, and adstock are the column variables. If you inspect the data you will see many missing adstock values, because the adstock variable is not defined for the first `N_lags` weeks in the data. To free memory, remove all missing values from `move` (`complete.cases`).

```
wide_data <- dcast(move, dma_code + store_code_uc + week_end ~ brand_name,
  value.var = c("quantity", "price", "promotion", "adstock"))
wide_data <- wide_data[complete.cases(wide_data), ]
```

```
head(wide_data)
```

Key: <dma_code, store_code_uc, week_end>

	dma_code	store_code_uc	week_end	quantity_comp	quantity_own	price_comp
	<int>	<int>	<Date>	<num>	<num>	<num>
1:	500	88153	2011-02-05	620	532	0.5834591
2:	500	88153	2011-02-12	870	210	0.5834591
3:	500	88153	2011-02-19	1432	238	0.5794119
4:	500	88153	2011-02-26	862	154	0.5860128
5:	500	88153	2011-03-05	1880	350	0.5860128
6:	500	88153	2011-04-09	762	210	0.5732326

	price_own	promotion_comp	promotion_own	adstock_comp	adstock_own
	<num>	<num>	<num>	<num>	<num>
1:	0.6667948	0	0	15.28524	17.20203
2:	0.6667948	0	0	16.10376	16.93872
3:	0.6667948	0	0	16.28562	15.65049
4:	0.6667948	0	0	16.55701	14.50100
5:	0.6667948	0	0	16.84442	11.60743
6:	0.6667948	0	0	16.35107	12.84814

3.6 Time fixed effects

Create an index for each month/year combination in the data using the following code:

```
wide_data[, week_end := as.Date(week_end, format = "%Y-%m-%d")]
wide_data[, month_index := 12 * (as.integer(format(week_end, "%Y")) - 2011) +
  as.integer(format(week_end, "%m"))]
head(wide_data[, .(week_end, month_index)]) #sanity check
```

	week_end	month_index
	<Date>	<num>
1:	2011-02-05	2
2:	2011-02-12	2
3:	2011-02-19	2
4:	2011-02-26	2
5:	2011-03-05	3
6:	2011-04-09	4

4 Data inspection

4.1 Time-series of advertising levels

We now take a look at the advertising data. First, pick a DMA. You can easily get a list of all DMA names and codes from the `stores` table. I picked "CHICAGO IL", which corresponds to `dma_code` 602. Then plot the time-series of weekly GRPs for your chosen market, separately for the own and competitor brand.

```
chosen_code = 506
move_filtered <- move[dma_code == chosen_code]
head(move_filtered)
```

Key: <brand_name, dma_code, week_end>

	brand_name	dma_code	week_end	store_code_uc	price	promotion	quantity
	<char>	<int>	<Date>	<int>	<num>	<num>	<num>
1:	comp	506	2010-01-02	12869	0.4862497	0.0000000	230
2:	comp	506	2010-01-02	18073	0.5466057	0.0000000	4660
3:	comp	506	2010-01-02	69801	0.5652718	0.3868085	3570
4:	comp	506	2010-01-02	73428	0.6423019	0.4711488	1036
5:	comp	506	2010-01-02	202709	0.5431721	0.0000000	2260
6:	comp	506	2010-01-02	208754	0.6484600	0.0000000	730

	dma_descr	grp_direct	grp_indirect	grp	adstock
	<char>	<num>	<num>	<num>	<num>
1: BOSTON (MANCHESTER)	MA-NH	219.0284	0	219.0284	NA
2: BOSTON (MANCHESTER)	MA-NH	219.0284	0	219.0284	NA
3: BOSTON (MANCHESTER)	MA-NH	219.0284	0	219.0284	NA
4: BOSTON (MANCHESTER)	MA-NH	219.0284	0	219.0284	NA
5: BOSTON (MANCHESTER)	MA-NH	219.0284	0	219.0284	NA
6: BOSTON (MANCHESTER)	MA-NH	219.0284	0	219.0284	NA

```
# Aggregate grp by week_end and brand
move_filtered <- move_filtered[, .(
  grp = mean(grp, na.rm = TRUE)
), by = .(week_end, brand_name)]
```


Note: I suggest you create a facet plot to display the time-series of GRPs for the two brands. Use the `facet_grid` or `facet_wrap` layer as explained in the `ggplot2` guide (see “More on facetting”).

```
library(ggplot2)

# Facet plot
facet_labels <- c("own" = "Own GRP Over Time",
                  "comp" = "Competitor GRP Over Time")

ggplot(move_filtered, aes(x = week_end, y = grp)) +
  geom_line(color = "blue") +
  facet_wrap(~ brand_name, ncol = 1, scales = "free_y",
             labeller = labeller(brand_name = facet_labels)) +
  labs(
    title = paste("Weekly GRP for DMA", chosen_code),
    x = "Week End",
    y = "GRP"
  ) +
  theme_minimal()
```

Weekly GRP for DMA 506



4.2 Overall advertising variation

Create a new variable **at the DMA-level**, `normalized_grp`, defined as $100 \times \text{grp} / \text{mean}(\text{grp})$. This variable captures the percentage deviation of the GRP observations relative to the DMA-level mean of advertising. Plot a histogram of `normalized_grp`.

```
mean_grp_dma_brand <- move[, .(mean_grp = mean(grp, na.rm = TRUE)), by = .(dma_code, brand_name)]

move <- merge(move, mean_grp_dma_brand,
              by = c("dma_code", "brand_name"), all.x = TRUE)
move[, normalized_grp := 100 * grp / mean_grp]


facet_labels <- c("own" = "Own GRP Ratio by DMA",
                  "comp" = "Competitor GRP Ratio by DMA")

# Set limits to exclude outliers
lower_limit <- 0
upper_limit <- 300

# Plot the histogram using 'move' without modifying it
ggplot(move, aes(x = normalized_grp)) +
  geom_histogram(binwidth = 10, fill = "green", color = "black",
                alpha = 0.7, boundary = 0) +
  facet_wrap(~ brand_name, ncol = 1, scales = "free_y",
            labeller = as_labeller(facet_labels)) +
  scale_x_continuous(limits = c(lower_limit, upper_limit)) + # Or use coord_cartesian()
  labs(
    title = "Histogram of Normalized GRP by Brand (Limited Range)",
    x = "Normalized GRP (%)",
    y = "Frequency"
  ) +
  theme_minimal()
```

Warning: Removed 23750 rows containing non-finite outside the scale range (``stat_bin()``).

Histogram of Normalized GRP by Brand (Limited)



Note: To visualize the data you should use the `scale_x_continuous` layer to set the axis limits. This data set is one of many examples where some extreme outliers distort the graph.

5 Advertising effect estimation

Estimate the following specifications:

1. Base specification that uses the log of `1+quantity` as output and the log of prices (own and competitor) and promotions as inputs. Control for store and month/year fixed effects.
2. Add the `adstock` (own and competitor) to specification 1.
3. Like specification 2., but not controlling for time fixed effects.

Combine the results using `etable` and comment on the results.

```
head(wide_data)
```

Key: <dma_code, store_code_uc>

	dma_code	store_code_uc	week_end	quantity_comp	quantity_own	price_comp
	<int>	<int>	<Date>	<num>	<num>	<num>
1:	500	88153	2011-02-05	620	532	0.5834591
2:	500	88153	2011-02-12	870	210	0.5834591
3:	500	88153	2011-02-19	1432	238	0.5794119
4:	500	88153	2011-02-26	862	154	0.5860128
5:	500	88153	2011-03-05	1880	350	0.5860128
6:	500	88153	2011-04-09	762	210	0.5732326

	price_own	promotion_comp	promotion_own	adstock_comp	adstock_own	month_index
	<num>	<num>	<num>	<num>	<num>	<num>
1:	0.6667948	0	0	15.28524	17.20203	2
2:	0.6667948	0	0	16.10376	16.93872	2
3:	0.6667948	0	0	16.28562	15.65049	2
4:	0.6667948	0	0	16.55701	14.50100	2
5:	0.6667948	0	0	16.84442	11.60743	3
6:	0.6667948	0	0	16.35107	12.84814	4

```
library(fixest)
```

```
basic_model <- feols(  
  log(1 + quantity_own) ~ log(price_own) + log(price_comp) + promotion_own +  
  promotion_comp |  
  store_code_uc + month_index,  
  data = wide_data  
)
```

```
model_with_adstock <- feols(  
  log(1 + quantity_own) ~ log(price_own) + log(price_comp) + promotion_own +  
  promotion_comp + adstock_own + adstock_comp |  
  store_code_uc + month_index,  
  data = wide_data  
)
```

```
model_with_adstock_no_time <- feols(  
  log(1 + quantity_own) ~ log(price_own) + log(price_comp) + promotion_own +  
  promotion_comp + adstock_own + adstock_comp |  
  store_code_uc,  
  data = wide_data  
)
```

```
etable(basic_model, model_with_adstock, model_with_adstock_no_time)
```

	basic_model	model_with_adstock	model_with_adstoc...
Dependent Var.: log(1+quantity_own)	log(1+quantity_own)	log(1+quantity_own)	log(1+quantity_own)
log(price_own)	-2.116*** (0.0819)	-2.113*** (0.0817)	-1.913*** (0.0770)
log(price_comp)	0.1436* (0.0579)	0.1466* (0.0578)	-0.4477*** (0.0787)
promotion_own	0.8782*** (0.0116)	0.8773*** (0.0116)	0.9286*** (0.0109)
promotion_comp	0.0346** (0.0113)	0.0349** (0.0113)	0.0075 (0.0147)
adstock_own		0.0307*** (0.0017)	-0.0182*** (0.0008)
adstock_comp		-0.0036*** (0.0007)	0.0113*** (0.0004)
Fixed-Effects:	-----	-----	-----
store_code_uc	Yes	Yes	Yes
month_index	Yes	Yes	No
S.E.: Clustered	by: store_code_uc	by: store_code_uc	by: store_code_uc
Observations	2,005,858	2,005,858	2,005,858
R2	0.64314	0.64321	0.63844
Within R2	0.07012	0.07030	0.07444

Signif. codes:	0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1		

The results from the basic model shows predominantly expected results. We see a negative correlation between own price and quantity sold, with a 2.1% decrease in quantity with each percentage point increase in the price of own products. There is a slight positive correlation between competitor price and own quantity sold, with own quantity going up by 0.14% for each percentage point competitors raise their prices. We see a positive correlation between own promotions and quantity, with the dummy variable having a coefficient of 0.8782. Slightly surprisingly, we also see a positive coefficient associated with competitor promotions (0.0346). This could be explained by a phenomenon where competitors' promotions draw some additional attention to own products as well, or causing customers to realize their need for a product within the category. This category spill-over means that there is a boost in sales of own products when competitors run a promotion on their own products.

The results from the adstock model with time fixed-effects show similar results to the basic model. It shows a slightly weaker negative correlation between own price and quantity sold, and a slightly stronger positive correlation between competitors' pricing and own quantity sold. The effects of running promotions is also similar to the basic model, with results showing a large positive effect of own promotions on quantities sold, and a small, but still positive effect of competitors running promotions. The adstock model also includes on the adstock, or the weighted average of current and past advertising, and these coefficients show the effects running an advertising campaign on the quantities sold. We can see that there is a slight positive correlation in own products sold when running an advertising campaign, while during and after a competitor's advertising, we see a decrease in the quantity of our own products sold.

In the model where we do not control for time fixed effects, our results deviate significantly from those in the previous models. While the coefficient on own price is still negative, indicating that an increase in the price of our products leads to a decrease in the number of units sold, the effect of competitors' price also has a negative coefficient. This would mean that an increase in a competing product's price lead to a decrease in sales for our own product as well. This is likely due to the model possibly overestimating the cross-price elasticity when time fixed effects are excluded. The effects of running promotions does not vary much in this model compared to the previous ones, though we do see that the coefficient on competitors' promotions is not statistically significant, indicating that the effect of a competitor running a promotion does not significantly change the sales metrics of our own product. The effect of adstock in this model is the reverse of what we observed in the model containing time fixed effects. Our results show a negative correlation between own advertising efforts and the quantity of own products sold, while it indicates a positive correlation between competitors' ad campaigns and our own products. This logically does not make sense and highlights the importance of including fixed effects into the regression model whenever possible.

Controlling for time fixed effects results in a higher coefficient in the effect of own adstock on the quantity

of own product sold when compared to the model without time fixed effects. From this, we can conclude that the company is focusing on advertising in markets where they believe they are not as strong as their competitors. The effect of competitor adstock, however, increases in the model without time fixed effects, which is in line with what we would expect when we do not control for time effects when the competition is focusing its advertising in markets where they already have a strong presence.